

# Knights Archers Zombies

Group number:

ST4406-G5

Group member:

ZhangJinwei

Zhu Hengyue

Zhang Ce



Related concepts



Key algorithm



Project background



Result display



Realization ideas  
and frame



Possible Improvement




Related concepts

- Agent: Agent Contains a policy and a set of internal values. It is the object of our learning and can make its own behavior.
- Environment: Receive action, generate state and reward, including a reward system, which may be random.
- Policy: The most basic function is to receive the Request, and then to provide the corresponding command. The command is the specific processing request.

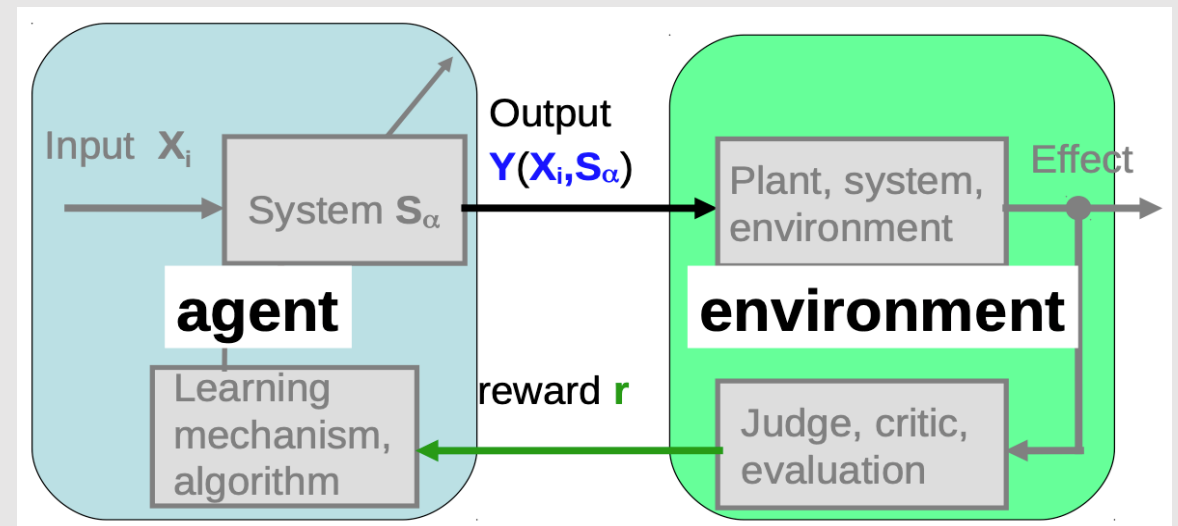


Gym is a toolkit for developing and comparing reinforcement learning algorithms. It supports teaching agents everything from walking to playing games like Pong or Pinball.



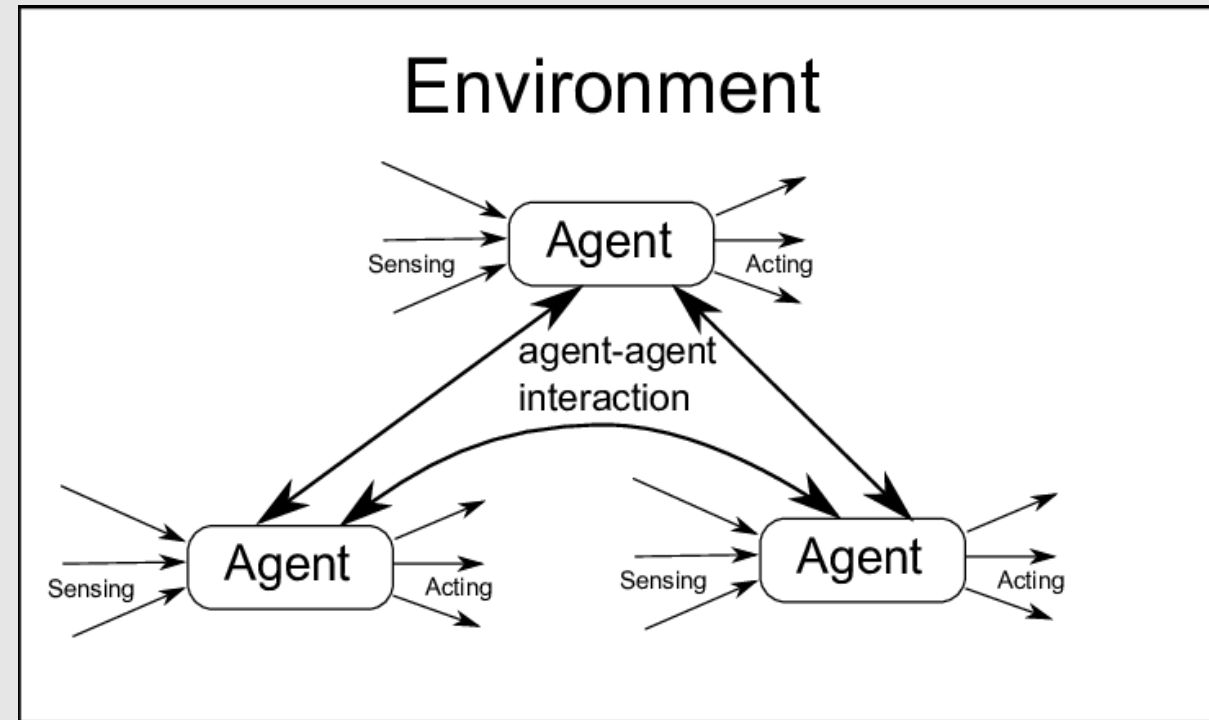
# Reinforcement Learning

Reinforcement stems from using machine learning to optimally control an agent in an environment. It works by learning a policy, a function that maps an observation obtained from its environment to an action. Policy functions are typically deep neural networks, which gives rise to the name “deep reinforcement learning.”



# Multiagent Reinforcement

In general it's the same as single agent reinforcement learning, where each agent is trying to learn its own policy to optimize its own reward. Using a central policy for all agents is possible, but multiple agents would have to communicate with a central server to compute their actions (which is problematic in most real world scenarios), so in practice decentralized multi-agent reinforcement learning is used.





Project background



# Rules introduction

- Zombies walk from the top border of the screen down to the bottom border in unpredictable paths.
- We control movable knights and archers. A knight is rewarded 1 point when its mace hits and kills a zombie. An archer is rewarded 1 point when one of their arrows hits and kills a zombie.
- The game ends when all agents die or a zombie reaches the bottom screen border.

# Our Goal

The goal is for the knights and archers to learn how to kill more zombies before the game is over



Realization ideas and frame

# Realization Ideas

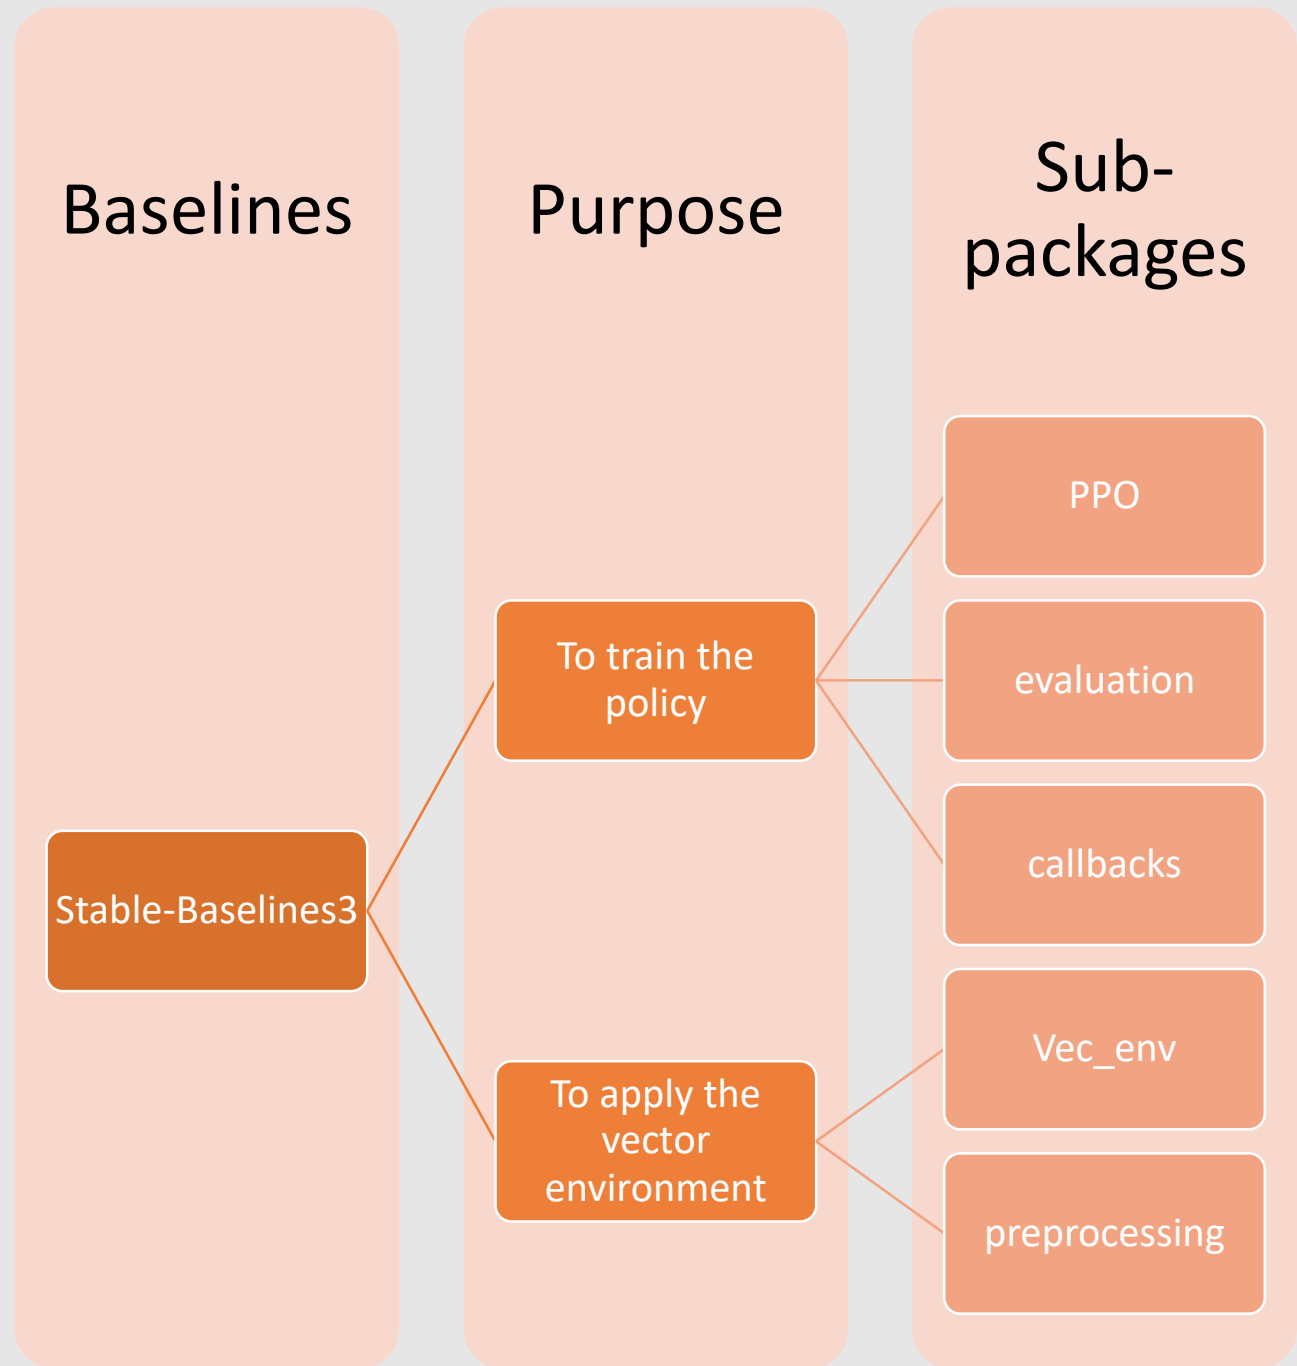
1

Use Stable-Baselines3 for evaluating policy and environment processing

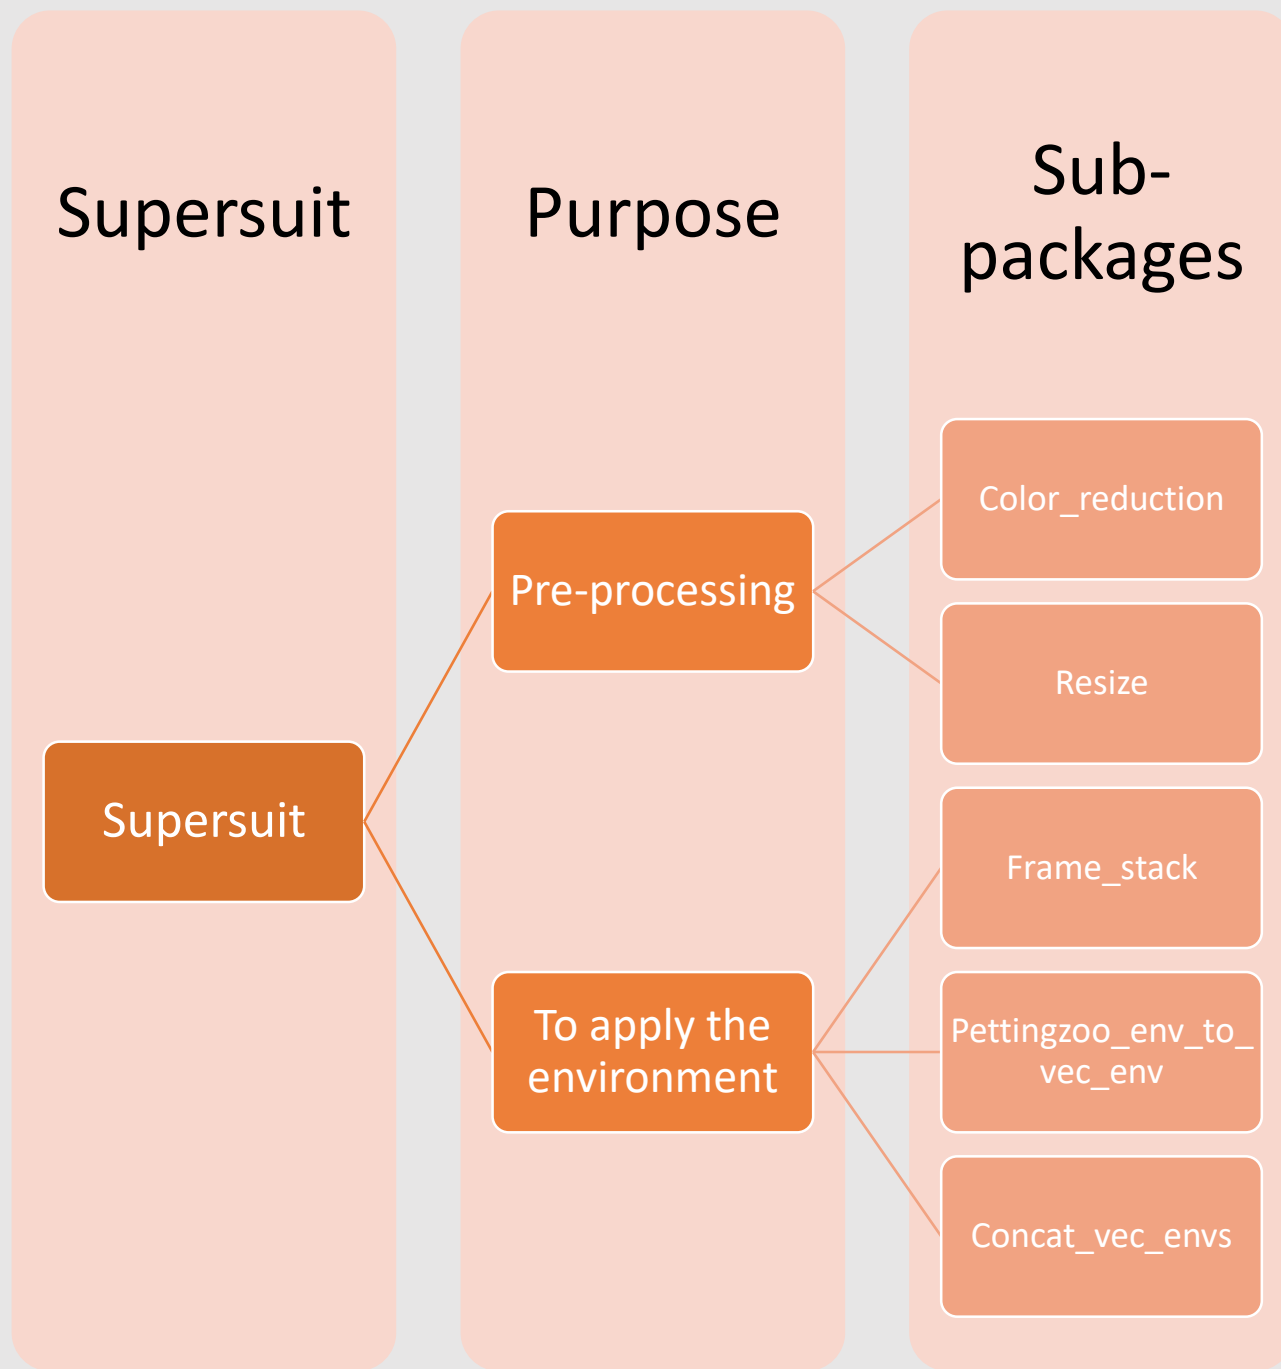
2

Use Supersuit for image simplifying and environment processing

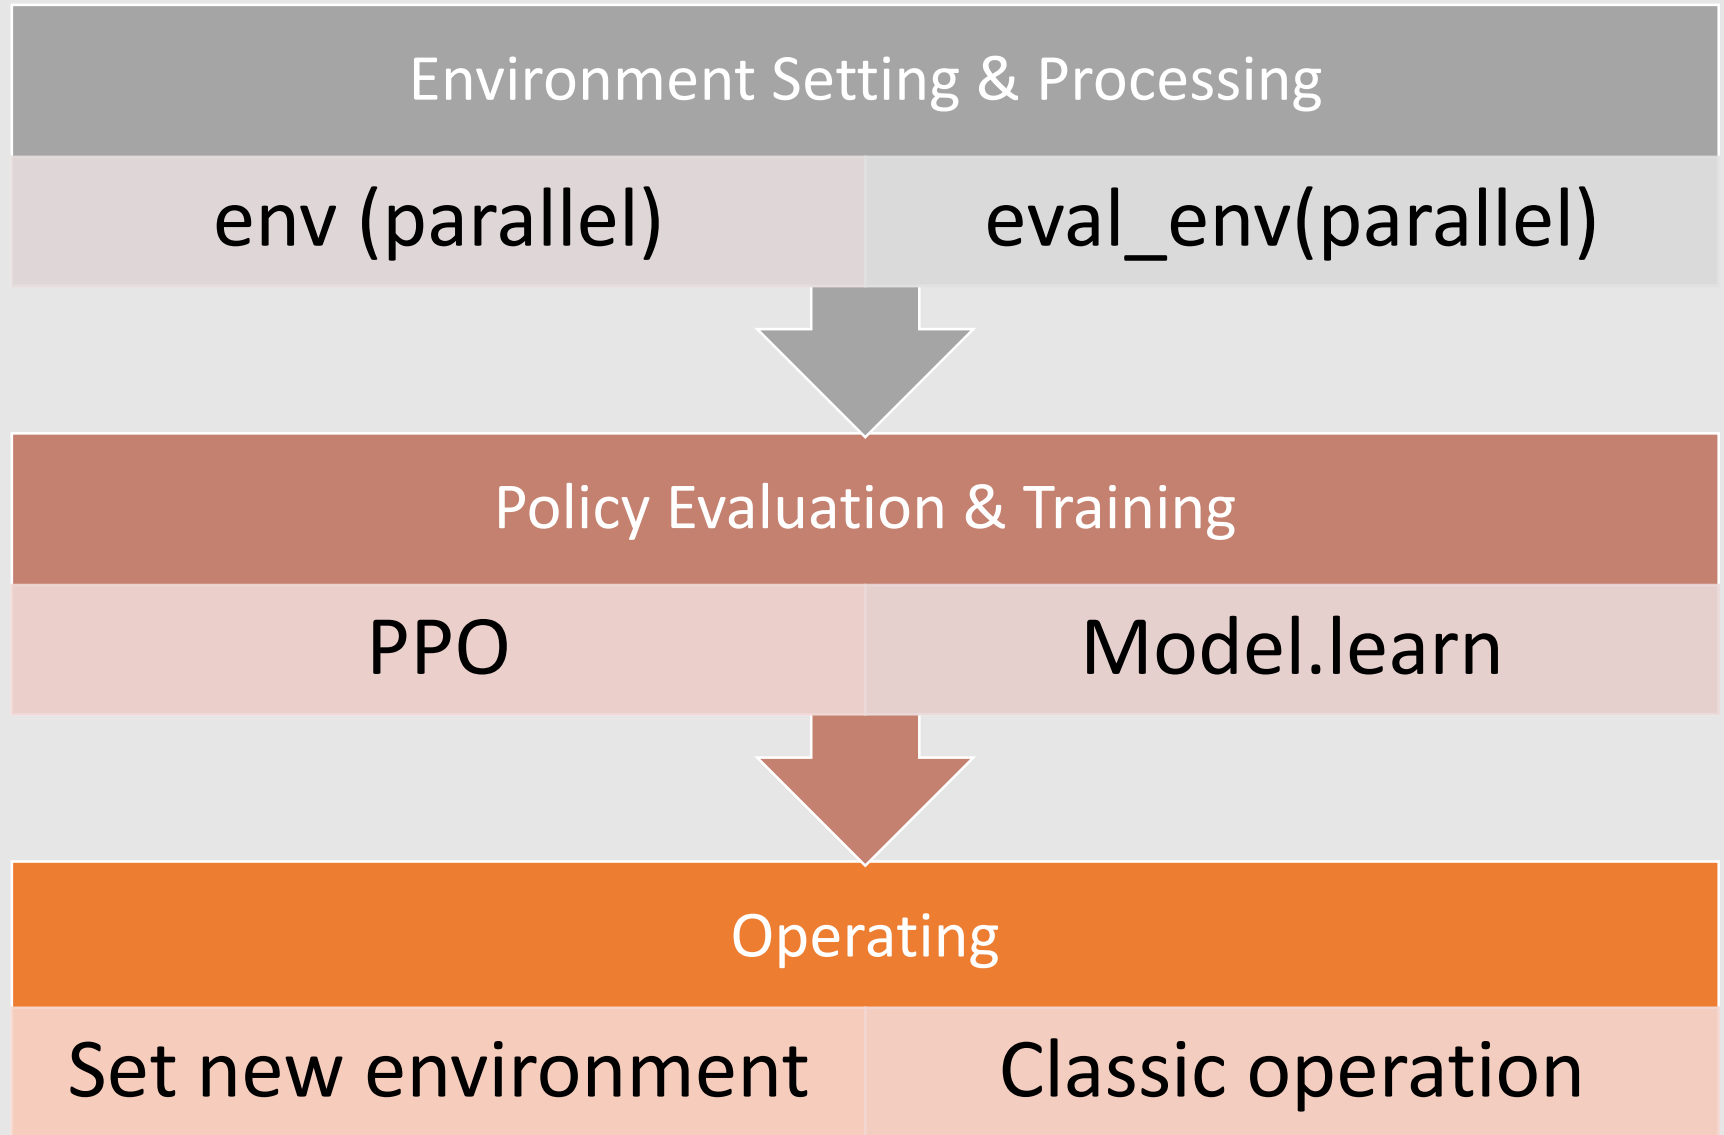
# Stable-Baselines3



# Supersuit



# Realization Frame





Key algorithm



# Proximal Policy Optimization

- Actor-Critic
- On policy
- Based on TRPO and PG
- Used the probability ratio instead if log probability
- Control the change of policy in each iteration
  - Add an adaptive KL penalty
  - Use CLIP, add an epsilon to control




$$L^{CLIP}(\theta) = \hat{E}_t[\min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1 - \varepsilon, 1 + \varepsilon)\hat{A}_t)]$$




- $\theta$  is the policy parameter
- $\hat{E}_t$  denotes the empirical expectation over timesteps
- $r_t$  is the ratio of the probability under the new and old policies, respectively
- $\hat{A}_t$  is the estimated advantage at time  $t$
- $\varepsilon$  is a hyperparameter, usually 0.1 or 0.2



Result display

	1	2	3	4	5	6	7	8	9	10
ep_rew_mean	0.875	1	1.09375	1.0652174	1.0925926	1.0735294	1.0897436	1.1304348	1.1	1.11
approx_kl		0.005459972	0.004744264	0.010000747	0.015741933	0.014996057	0.026188482	0.025564905	0.032294594	0.019883174
clip_fraction		0.000488	0.00713	0.00869	0.0324	0.0711	0.122	0.12	0.148	0.0807
entropy_loss		-1.79	-1.79	-1.78	-1.79	-1.77	-1.74	-1.75	-1.76	-1.77
loss		-0.163	-0.177	-0.165	-0.181	-0.182	-0.201	-0.181	-0.185	-0.182
policy_gradient_loss		-0.00116	-0.00304	-0.00157	-0.00644	-0.0147	-0.0171	-0.00964	-0.0132	-0.0158
value_loss		0.0801	0.0589	0.0609	0.0439	0.0415	0.0447	0.0684	0.0555	0.0471
	11	12	13	14	15	16	17	18	19	20
ep_rew_mean	1.12	1.02	0.9	0.89	0.75	0.66	0.59	0.6	0.64	0.71
approx_kl	0.025	0.028212432	0.03647689	0.03683254	0.038915418	0.040916942	0.028782278	0.036273815	0.053855795	0.037858583
clip_fraction	0.125	0.103	0.199	0.173	0.182	0.179	0.174	0.139	0.204	0.128
entropy_loss	-1.77	-1.76	-1.77	-1.75	-1.74	-1.74	-1.75	-1.74	-1.75	-1.75
loss	-0.17	-0.197	-0.183	-0.199	-0.201	-0.174	-0.199	-0.198	-0.17	-0.19
policy_gradient_loss	-0.01	-0.0182	-0.0151	-0.0265	-0.021	-0.0202	-0.016	-0.0122	-0.0184	-0.0172
value_loss	0.048	0.0447	0.0278	0.0128	0.0323	0.0221	0.019	0.0365	0.0369	0.0695







Possible Improvement

# Possible Improvement

Maybe 2 systems with different  
policy

# Two System

System 1:

Use the simple policy, take action by  
observetion

System 2:

Use PPO to train the model and take the  
optimized policy



# Reference

<https://towardsdatascience.com/multi-agent-deep-reinforcement-learning-in-15-lines-of-code-using-pettingzoo-e0b963c0820b>

[https://github.com/jkterry1/rl\\_scratch/blob/a5476ce2332e243dafd5bd804c3bac5e7ae176f2/test\\_evaluation.py](https://github.com/jkterry1/rl_scratch/blob/a5476ce2332e243dafd5bd804c3bac5e7ae176f2/test_evaluation.py)

<https://github.com/gml16/yare-rl/blob/6d24a596eb870f54d13898dc76c5da2489135190/yare-rl/train.py>

**THANK YOU**